

Assignment-4

Name: Prakhar Agarwal

Reg No.: 21BIT0034

Email ID: prakhar.agarwal2021@vitstudent.ac.in

TASK 1

```
import pandas as pd
```

```
df=pd.read_csv('/content/winequality-red.csv')  
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 1599 entries, 0 to 1598  
Data columns (total 12 columns):  
#   Column                Non-Null Count  Dtype  
---  ---  
0   fixed acidity          1599 non-null   float64  
1   volatile acidity       1599 non-null   float64  
2   citric acid            1599 non-null   float64  
3   residual sugar         1599 non-null   float64  
4   chlorides              1599 non-null   float64  
5   free sulfur dioxide    1599 non-null   float64  
6   total sulfur dioxide   1599 non-null   float64  
7   density                1599 non-null   float64  
8   pH                    1599 non-null   float64  
9   sulphates              1599 non-null   float64  
10  alcohol                1599 non-null   float64  
11  quality                1599 non-null   int64  
dtypes: float64(11), int64(1)  
memory usage: 150.0 KB
```

```
df.shape
```

```
(1599, 12)
```

TASK 2

Univariate

```
import pandas as pd  
import numpy as np  
import matplotlib.pyplot as plt  
import seaborn as sns
```


```
sns.displot(df.pH)
```

```
<seaborn.axisgrid.FacetGrid at 0x7df31745cc40>
```



```
plt.pie(df.pH.value_counts())  
plt.title('WATER')  
plt.show()
```


WATER



Bivariate


```
sns.lineplot(x = df.pH,y=df.density)
```

```
<Axes: xlabel='pH', ylabel='density'>
```



```
sns.scatterplot(x = df.chlorides,y=df.sulphates)
```


<Axes: xlabel='chlorides', ylabel='sulphates'>



Multivariate


| |
sns.pairplot(df)

```
<seaborn.axisgrid.PairGrid at 0x7df2fd3d3910>
```



```
sns.heatmap(df.corr(),annot=True)
```

```
<Axes: >
```



TASK 3



```
from sklearn.linear_model import LinearRegression
```

```
df.isnull().sum()
```

```
fixed acidity      0
volatile acidity   0
citric acid        0
residual sugar     0
chlorides          0
free sulfur dioxide 0
total sulfur dioxide 0
density            0
pH                 0
sulphates          0
alcohol            0
quality            0
dtype: int64
```

```
df.value_counts()
```

```
fixed acidity  volatile acidity  citric acid  residual sugar  chlorides  free sulfur dioxide  total sulfur dioxide  density  pH
sulphates  alcohol  quality
7.5          0.510          0.02          1.7          0.084          13.0          31.0          0.99538  3.36
0.54         10.5         6          4          0.077          18.0          34.0          0.99480  3.39
6.7          0.460          0.24          1.7          0.074          24.0          44.0          0.99534  3.40
0.60         10.6         6          4          0.076          12.0          20.0          0.99546  3.29
7.2          0.360          0.46          2.1          0.089          17.0          40.0          0.99803  3.29
0.85         11.0         7          4          0.13          2.0          20.0          0.99546  3.29
0.54         10.1         5          4          0.10          2.9          40.0          0.99803  3.29
8.3          0.650          0.10          2.9          0.062          6.0          12.0          0.99495  3.44
0.55         9.5         5          3          0.32          2.7          25.0          0.99628  3.24
```

```

0.44      10.4      5      1
          0.31      1.6      0.080      15.0      42.0      0.99780  3.31
0.64      9.0      5      1
          0.410      0.15      3.7      0.104      29.0      94.0      0.99786  3.14
0.58      9.1      5      1
15.9      0.360      0.65      7.5      0.096      22.0      71.0      0.99760  2.98
0.84      14.9      5      1
Length: 1359, dtype: int64

```

```
df.nunique()
```

```

fixed acidity      96
volatile acidity   143
citric acid        80
residual sugar     91
chlorides          153
free sulfur dioxide 60
total sulfur dioxide 144
density           436
pH                89
sulphates         96
alcohol           65
quality           6
dtype: int64

```

```

y = df['pH']
y.head()

```

```

0    3.51
1    3.20
2    3.26
3    3.16
4    3.51
Name: pH, dtype: float64

```

```

X= df.drop(columns = ['quality'],axis =1)
X.head()

```

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcohol
0	7.4	0.70	0.00	1.9	0.076	11.0	34.0	0.9978	3.51	0.56	9.4
1	7.8	0.88	0.00	2.6	0.098	25.0	67.0	0.9968	3.20	0.68	9.8
2	7.8	0.76	0.04	2.3	0.092	15.0	54.0	0.9970	3.26	0.65	9.8
3	11.2	0.28	0.56	1.9	0.075	17.0	60.0	0.9980	3.16	0.58	9.8
4	7.4	0.70	0.00	1.9	0.076	11.0	34.0	0.9978	3.51	0.56	9.4

```

from sklearn.preprocessing import StandardScaler
scale =StandardScaler()

```

```

X_scaled =pd.DataFrame(scale.fit_transform(X),columns = X.columns)
X_scaled.head()

```

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphate
0	-0.528360	0.961877	-1.391472	-0.453218	-0.243707	-0.466193	-0.379133	0.558274	1.288643	-0.57920
1	-0.298547	1.967442	-1.391472	0.043416	0.223875	0.872638	0.624363	0.028261	-0.719933	0.12895
2	-0.298547	1.297065	-1.186070	-0.169427	0.096353	-0.083669	0.229047	0.134264	-0.331177	-0.04808
3	1.654856	-1.384443	1.484154	-0.453218	-0.264960	0.107592	0.411500	0.664277	-0.979104	-0.46118
4	-0.528360	0.961877	-1.391472	-0.453218	-0.243707	-0.466193	-0.379133	0.558274	1.288643	-0.57920

```

from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test = train_test_split(X_scaled,y,test_size =0.2,random_state =0)

```

```

x_train.shape

(1279, 11)

```

```
x_test.shape
```

```
(320, 11)
```

```
x_test.head()
```

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulph
1109	1.425044	-0.323013	0.816598	-0.311323	1.775397	1.063900	0.593954	0.770280	-0.914312	0.60
1032	-0.126188	1.632254	-1.391472	1.107633	0.160114	-1.039977	-0.987312	0.950485	0.316751	-0.75
1002	0.448342	-1.328579	0.303093	-0.346797	-0.520005	-0.274931	-0.591995	-0.840962	-0.331177	1.07
487	1.080326	0.654620	0.457144	-0.524166	-0.732542	-1.039977	-0.987312	0.770280	-0.914312	-1.40
979	2.229387	-0.434742	1.124700	-0.807957	-0.264960	-1.231239	-1.230584	0.081262	-1.173483	-0.16

```
y_test
```

```
1109    3.17
1032    3.36
1002    3.26
487     3.17
979     3.13
...
794     3.17
813     3.44
1322    3.18
704     3.29
1023    3.27
Name: pH, Length: 320, dtype: float64
```

```
from sklearn.linear_model import LinearRegression
lr = LinearRegression()
```

```
lr.fit(x_train,y_train)
```

```
LinearRegression
LinearRegression()
```

```
y_predict = lr.predict(x_test)
y_predict
```

```
array([3.17, 3.36, 3.26, 3.17, 3.13, 3.13, 3.12, 3.47, 3.3 , 3.27, 3.2 ,
       3.14, 3.29, 3.74, 3.5 , 3.34, 3.33, 3.15, 3.55, 3.38, 3.35, 3.27,
       3.36, 3.3 , 3.57, 3.27, 3.23, 3.35, 3.2 , 3.9 , 3.14, 3.31, 3.53,
       3.5 , 3.52, 3.26, 3.46, 3.69, 3.62, 3.35, 3.33, 3.37, 3.49, 3.02,
       3.3 , 3.19, 3.16, 3.3 , 3.58, 3.36, 3.41, 3.38, 3.38, 3.3 , 3.32,
       3.35, 3.3 , 3.09, 3.23, 3.27, 3.68, 3.32, 3.32, 3.56, 3.51, 3.42,
       3.29, 2.92, 3.19, 3.07, 3.3 , 3.04, 3.38, 3.17, 3.34, 3.58, 3.46,
       3.45, 3.35, 3.3 , 3.1 , 3.39, 3.51, 3.14, 3.36, 3.4 , 3.44, 3.25,
       3.42, 3.44, 3.26, 3.22, 3.62, 3.35, 3.4 , 3.3 , 3.15, 3.19, 3.26,
       3.04, 3.23, 3.38, 3.29, 3.34, 3.46, 3.39, 3.46, 3.41, 3.42, 3.23,
       3.56, 3.1 , 3.23, 3.29, 3.33, 3.37, 3.25, 3.27, 3.45, 3.36, 3.23,
       2.92, 3.03, 3.36, 3.27, 3.3 , 3.23, 3.13, 3.36, 3.06, 3.4 , 3.35,
       3.32, 3.22, 3.28, 3.29, 3.06, 3. , 2.88, 3.36, 3.39, 3.19, 3.16,
       3.67, 3.17, 3.3 , 3.23, 3.33, 3.18, 3.34, 3.42, 3.39, 3.41, 3.29,
       3.61, 3.46, 3.25, 3.28, 3.44, 3.38, 3.48, 3.42, 3.32, 3.5 , 3.14,
       3.25, 3.54, 3.41, 3.18, 3.27, 3.1 , 3.35, 3.48, 3.15, 3.52, 3.21,
       3.18, 3.34, 3.71, 3.16, 3.32, 3.38, 3.41, 3.2 , 3.39, 3.2 , 3.26,
       3.29, 3.25, 3.42, 3.33, 2.94, 3.18, 3.29, 3.46, 3.28, 3.2 , 3.39,
       3.28, 3.16, 3.19, 3.41, 3.26, 3.31, 3.26, 3.59, 3.24, 3.38, 3.36,
       3.42, 3.29, 3.25, 3.11, 3.18, 3.49, 3.47, 3.04, 3.39, 3.38, 3.28,
       3.51, 3.16, 3.2 , 3.03, 3.36, 3.12, 3.22, 3.56, 3.16, 3.39, 3.53,
       3.57, 3.35, 3.23, 3.13, 3.42, 3.07, 3.33, 3.36, 3.22, 3.48, 3.4 ,
       3.36, 3.39, 3.24, 3.26, 3.29, 3.29, 3. , 3.39, 3.19, 3.44, 3.23,
       3.26, 3.4 , 3.36, 3.21, 3.11, 3.48, 3.07, 3.44, 3.68, 3.14, 3.08,
       3.34, 3.36, 3.6 , 3.28, 3.28, 3.33, 3.14, 3.17, 3.24, 3.9 , 3.22,
       3.39, 3.14, 3.21, 3.39, 3.05, 3.14, 3.27, 3.17, 3.26, 3.26, 3.55,
       3.36, 3.27, 3.34, 3.15, 3.49, 3.5 , 3.41, 3.19, 3.26, 3.38, 3.4 ,
       3.29, 3.35, 3.15, 3.44, 3.56, 3.12, 3.31, 3.35, 3.61, 3.45, 3.48,
       3.37, 3.39, 3.36, 3.42, 3.66, 3.19, 3.2 , 3.17, 3.44, 3.18, 3.29,
       3.27])
```

```

y_predict1 =lr.predict(x_train)
y_predict1

array([3.39, 3.13, 3.26, ..., 3.29, 3.3 , 3.25])

profit =pd.DataFrame({'Actual_pH':y_test,'Predicted_pH':y_predict})
profit

```

	Actual_pH	Predicted_pH
1109	3.17	3.17
1032	3.36	3.36
1002	3.26	3.26
487	3.17	3.17
979	3.13	3.13
...
794	3.17	3.17
813	3.44	3.44
1322	3.18	3.18
704	3.29	3.29
1023	3.27	3.27

320 rows × 2 columns

Task 4

```

from sklearn import metrics

print(metrics.r2_score(y_test,y_predict))

1.0

print(metrics.r2_score(y_train,y_predict1))

1.0

print(metrics.mean_squared_error(y_test,y_predict))

3.457429436163966e-31

print(np.sqrt(metrics.mean_squared_error(y_test,y_predict)))

```

5.879991017139368e-16

+ Code

+ Text

TASK 5

```
df.head()
```

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulp
0	7.4	0.70	0.00	1.9	0.076	11.0	34.0	0.9978	3.51	
1	7.8	0.88	0.00	2.6	0.098	25.0	67.0	0.9968	3.20	
2	7.8	0.76	0.04	2.3	0.092	15.0	54.0	0.9970	3.26	
3	11.2	0.28	0.56	1.9	0.075	17.0	60.0	0.9980	3.16	

```
lr.predict([[7.4, 0.70, 0.00, 1.9, 0.076, 11.0, 34.0, 0.9978, 3.51, 0.56, 9.4]])/100000
```

```


/usr/local/lib/python3.10/dist-packages/sklearn/base.py:439: UserWarning: X does not have valid feature names, but LinearRegression was
warnings.warn(

```

```
array([3.85284021e-05])
```

```
lr.predict([[7.8, 0.88, 0.00, 2.6, 0.098, 25.0, 67.0, 0.9968, 3.20, 0.68, 9.8]])/100000
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/base.py:439: UserWarning: X does not have valid feature names, but LinearRegression was  
warnings.warn(  
array([3.80499538e-05])
```



```
lr.predict([[7.4, 0.70, 0.00, 1.9, 0.076, 11.0, 34.0, 0.9978, 3.51, 0.56, 9.4]])/100000
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/base.py:439: UserWarning: X does not have valid feature names, but LinearRegression was  
warnings.warn(  
array([3.85284021e-05])
```

